Modular Inverse Reinforcement Learning on Human Motion

Shun Zhang
Department of Computer Science
University of Texas at Austin
Austin, TX 78712
menie482@cs.utexas.edu

Matthew Tong
Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
mhtong@gmail.com

Mary Hayhoe
Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
hayhoe@utexas.edu

Dana Ballard
Department of Computer Science
University of Texas at Austin
Austin, TX 78712
dana@cs.utexas.edu

Abstract

1 Introduction

Human is able to learn accomplishing complicated tasks much faster than machines can do. However, most of the tasks can be decomposed into subtasks. A human may already have the capacity to accomplish the subtasks. He simply transfers the knowledge of the subtasks and integrate them for the objective task. For example, when a human learns to drive, he has much prior domain knowledge to help him with this objective task. He integrates his skills of controlling the velocity, avoiding other vehicles, navigating, responding to traffic lights, and so on.

Similar ideas are also adopted in the learning literature. Reinforcement learning suffers from the curse of dimensionality. It is inefficient for a learning agent to learn a complicated task from scratch. So learning algorithms including hierarchical reinforcement learning [1], modular reinforcement learning [5] are proposed. They also decompose the objective task into subtasks. The learning algorithms first learn the subtasks independently, then learn how to combine these subtasks.

A natural question to ask is how human integrates skills for subtasks, and whether we can use the same method for integrating subtasks in reinforcement learning problems. In this paper, we analyze human's behavior of accomplishing a task composite of various subtasks. We collected human motion data and try to interpret the behavior using a way of inverse reinforcement learning. With the best of our knowledge, this approach is novel in the literature.

This paper is organized as follows. Section 2 introduces the domain of the composite task that we collected human data. Section 3 describes the main algorithm, modular inverse reinforcement learning. We report our experiment results in Section 4, and conclude in Section 5.

2 Multi-objective Sidewalk Domain

Figure 1: The task of collecting targets, avoiding obstacles and following the path. [4]

Figure 1 shows the domain that we use in this paper. The pink cylinders represent obstacles, that the player would be punished if running into. The blue cylinders represent targets, that the player would be rewarded if collected. There is also a gray path on the ground that the player can follow. So, naturally, this domain has three subtasks, 1) following the path, 2) collecting targets, 3) avoiding obstacles. This is an experiment used in the literature to evaluate modular reinforcement learning [4].

We will evaluate four different tasks. **Task 1**, following the path only, and ignoring other objects. **Task 2**, following the path, while avoid the obstacles. **Task 3**, following the path, while attain targets when possible. **Task 4**, following the path, collecting the targets and avoiding obstacles simultaneously. We conducted experiments to ask volunteers to accomplish different tasks and recorded their trajectories. The human data are collected by the Center for Perceptual Systems in University of Texas at Austin.

If the player observes the distance and angle to an object, as shown in Figure 1 B), he is expected to know the optimal action to avoid or pursue it. This decision is also Markov. Therefore, the question would be, if we have trained modules for these three subtasks, could they be integrated to match human's behavior?

3 Modular Inverse Reinforcement Learning

We need some basic concepts in reinforcement learning to proceed. For a state, action pair, s and a, $Q(s, a)$ evaluates the utility of taking action a from state s. The *policy* of a state is the action with the maximum Q value[6]. So how to we determine the global policy from the modules, or subtasks? In the literature, work has been done to integrate the Q function [2] or compromise on policies directly [7].

Here, we assume that the global Q function is a weighted sum of all Q_i, where Q_i is the Q function for i-th module.

$$Q(s,a) = \sum_i w_i Q_i(s_i, a)$$

where w_i is the weight of the i-th sub-MDP. $w_i \geq 0, \sum_i w_i = 1$. s_i denotes the decomposition of s in the i-th module.

Different weights can yield different performance. Let w_1, w_2, w_3 be weights for the task of target collection, obstacle avoidance, and path following, respectively. Let w be the vector of (w_1, w_2, w_3). An agent with $w = (1, 0, 0)$ only collects targets, and one with $w = (0, 0.5, 0.5)$ may avoid the obstacles and follow the path.

To obtain the weights given the samples, we need to use the Inverse Modular Reinforcement Learning technique [4]. We use a maximum likelihood method here.

$$\max_w \prod_t \frac{e^{\eta Q(s^{(t)}, a^{(t)})}}{\sum_b e^{\eta Q(s^{(t)}, b)}} \tag{1}$$

where $s^{(t)}$ is the state at time t, and $a^{(t)}$ is the action at time t, which are both from samples. $Q(s,a) = \sum_i w_i Q_i(s_i, a)$, as defined before. η is a hyperparameter that determines the consistency of human's behavior. The larger η is, the algorithm is more likely to overfit the data.

The intuition of Equation 1 is that if an action is observed from the sample, then the Q value of taking that action should be larger compared to Q values of taking other actions.

4 Experiments

We train the modules first to get Q_i before running the inverse reinforcement learning algorithm. For each module, the agent only considers the closest target and the closest obstacle. For the path module, the path is defined as segments of waypoints, so the closest waypoint is considered. The agent takes the distance and angle to the closest objects as the state representation.

To make our weights represent the significance of the modules, we normalize the sub-MDPs with the unit (positive or negative) rewards. The reward is 1 for collecting a target, -1 for running into an obstacle. We define the value function directly for the path module to have a path following performance, as it is tricky to give reward for such performance.

We make some constraints on our learning agent to make it walk like a human. We can find in the human trajectories that humans walk smoothly. They don't turn sharply. Our agent is only allowed to do three actions — going straight ahead, turning left with a small step, and turning right with a small step.


Figure 2: The trajectories of humans and our agent in four different tasks. From left to right are Task 1 to Task 4. The active modules should be (Path), (Obstacle + Path), (Target + Path), (Target + Obstacle + Path), respectively.

Average by Task	Num Targs Hit	Num Obst Hit
1	2.34	2.13
2	3.03	**0.13**
3	**10.19**	2.28
4	**9.88**	**0.03**

Table 1: Number of targets hit and number of obstacles hit of the humans.

Average by Task	Num Targs Hit	Num Obst Hit
1	1.25	1.62
2	3.62	**2.37**
3	**5.14**	3.14
4	**5.00**	**2.00**

Table 2: Number of targets hit and number of obstacles hit of the learning agent.

We report the results in Figure 2, same as Figure 1, the red circles are obstacles. The blue circles are targets. The gray lines are the path. The black lines are trajectories of human — each line represents one human trajectory. Using the weights derived from the algorithm, the trajectories of our agents are drawn in green color.

Although humans don't agree each other on how to avoid obstacles or collecting targets, our agent can figure out what the humans are doing, and perform a similar trajectory. Table 1 and Table 2 evaluate the performance by showing the number of targets hit and number of obstacles hit. Note that the numbers bold-ed are active modules in the corresponding task. We can observe that humans still do better than our agent in these tasks.


Figure 3: Heatmaps of the $\log$ of values of Equation 1 for different weights for the four tasks, respectively. The red zones indicate higher values. The upper two are Task 1 and 2. The lower two are Task 3 and 4.

In Figure 3, we show the $\log$ of values of Equation 1 for different weights. We can observe the centroids of red zones move for different tasks. It stays at the origin in Task 1, so none of target and obstacle modules are active. It moves away from the origin when a module is active.

From the heatmaps, we find the optimums exist for these tasks. The optimal weights are also consistent with the experiment context.

5   Conclusion and Future Work

We interpreted human behavior using inverse module reinforcement learning in this paper. We have some positive results, but the performance of our agent is still inferior than humans.

There are also some observations potential for the future work. First, weighted sum of Q functions is one way to combine multiple sub-MDPs. We also propose other ways including, for example, scheduling between different modules, with only one active at one time. This is also called skilled in the literature [3]. However, we adopt the weighted sum approach because this is more reasonable for human behavior. When a human tries to collect targets while avoiding obstacles, these two modules are expected to be both active. A scheduling approach may yield frequent oscillation between these two modules.

Second, we also assumes independency between modules. Correlation between modules doesn't impair our analysis in this paper. In Figure 3, we can tell that the target module and obstacle module tend to be negatively correlated from the shape of the red zones.

Lastly, weights may be dynamic and different from state to state. However, with such assumption, we need to learn a mapping from state to weights. In this case, the curse of dimensionality still exists, and inverse learning would be difficult.

References

[1] Thomas G Dietterich. Hierarchical reinforcement learning with the maxq value function decomposition. *J. Artif. Intell. Res.(JAIR)*, 13:227–303, 2000.

[2] Daphne Koller and Ronald Parr. Computing factored value functions for policies in structured mdps. In *IJCAI*, volume 99, pages 1332–1339, 1999.

[3] George Konidaris and Andre S Barreto. Skill discovery in continuous reinforcement learning domains using skill chaining. In *Advances in Neural Information Processing Systems*, pages 1015–1023, 2009.

[4] Constantin A Rothkopf and Dana H Ballard. Modular inverse reinforcement learning for visuomotor behavior. *Biological cybernetics*, 107(4):477–490, 2013.

[5] Nathan Sprague and Dana Ballard. Multiple-goal reinforcement learning with modular sarsa (0). In *IJCAI*, pages 1445–1447. Citeseer, 2003.

[6] Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*, volume 1. Cambridge Univ Press, 1998.

[7] Philip S Thomas and Andrew G Barto. Motor primitive discovery. In *ICDL-EPIROB*, pages 1–8, 2012.