
Modular Inverse Reinforcement Learning on Human Motion

Shun Zhang

Department of Computer Science
University of Texas at Austin
Austin, TX 78712
menie482@cs.utexas.edu

Matthew Tong

Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
mhtong@gmail.com

Mary Hayhoe

Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
hayhoe@utexas.edu

Dana Ballard

Department of Computer Science
University of Texas at Austin
Austin, TX 78712
dana@cs.utexas.edu

Abstract

Keywords: Reinforcement learning, human motion, inverse learning

1 Introduction

A large body of evidence from human and other primate studies suggests that they use reinforcement signals to learn tasks and reinforcement learning has proved to be a good model for this process. However humans are able to learn and carry out very complex tasks involving many different objectives.

How can they do this? The likelihood is that they use reinforcement in this context as well but most formal reinforcement learning models do not scale well with increasing complexity.

One promising possibility is that the complex task can be broken down into subtasks that are each learned separately [4, 3]. For certain task venues, this decomposition allows the behavior in the complex task to be specified as a weighted sum of individual sub-task priorities. This paper explores the use of such a decomposition to characterize complex behavior of subjects in a structure virtual reality environment. The methodology of [3] allows the task priorities to be estimated using the subjects' behavioral data.

Our experimental analysis shows that the modular reinforcement can be a surprisingly good model of behavior, predicting individual subjects' sub-task priorities in a way that explains their behavioral choices.

This paper is organized as follows. Section 2 introduces the domain of the composite task that we collected human data. Section 3 describes the main algorithm, modular inverse reinforcement learning. We report our experiment results in Section 4, and conclude in Section 5.

2 Multi-objective Domain



Figure 1: (Left) A volunteer with head tracker, virtual reality display, and body tracker. (Right) The environment the human can see through the display.

In Figure 1 shows the domain that we use in this paper. The red cubes represent obstacles, that the player would be punished if running into. The blue balls represent targets, that the player would be rewarded if collected. There is also a gray path on the ground that the player can follow. Naturally, this domain has three subtasks, 1) following the path, 2) collecting targets, 3) avoiding obstacles. This is an experiment used in the literature to evaluate modular reinforcement learning [3]. We ask our volunteers to have trackers and virtual reality displays equipped. He can see the environment through the display in front of his eyes. So he can walk as if he is walking in the virtual domain.

We will evaluate four different tasks. **Task 1**, following the path only, and ignoring other objects. **Task 2**, following the path, while avoid the obstacles. **Task 3**, following the path, while attain targets when possible. **Task 4**, following the path, collecting the targets and avoiding obstacles simultaneously. We conducted experiments to ask volunteers to accomplish different tasks and recorded their trajectories. The human data are collected by the Center for Perceptual Systems at University of Texas at Austin.

If the player observes the distance and angle to an object, he is expected to know the optimal action to avoid or pursue it. This decision is also Markov. Therefore, the question would be, if we have trained modules for these three subtasks, could they be integrated to match human's behavior?

3 Modular Inverse Reinforcement Learning

For a state, action pair, s and a , $Q(s, a)$ evaluates the utility of taking action a from state s . The *policy* of a state is the action with the maximum Q value [5]. So how to we determine the global policy from the modules, or subtasks? In the literature, work has been done to integrate the Q function [1] or compromise on policies directly [6].

Here, we assume that the global Q function is a weighted sum of all Q_i , where Q_i is the Q function for i-th module.

$$Q(s, a) = \sum_i w_i Q_i(s_i, a)$$

where w_i is the weight of the i-th sub-MDP. $w_i \geq 0$, $\sum_i w_i = 1$. s_i denotes the decomposition of s in the i-th module.

Different weights can yield different performance. Let w_1, w_2, w_3 be weights for the task of target collection, obstacle avoidance, and path following, respectively. Let w be the vector of (w_1, w_2, w_3) . An agent with $w = (1, 0, 0)$ only collects targets, and one with $w = (0, 0.5, 0.5)$ may avoid the obstacles and follow the path.

To obtain the weights given the samples, we need to use the Inverse Modular Reinforcement Learning technique [3]. We use a maximum likelihood method here.

$$\max_w \prod_t \frac{e^{\eta Q(s^{(t)}, a^{(t)})}}{\sum_b e^{\eta Q(s^{(t)}, b)}} \quad (1)$$

where $s^{(t)}$ is the state at time t , and $a^{(t)}$ is the action at time t , which are both from samples. $Q(s, a) = \sum_i w_i Q_i(s_i, a)$, as defined before. η is a hyperparameter that determines the consistency of human's behavior. The larger η is, the algorithm is more likely to overfit the data.

The intuition of Equation 1 is that if an action is observed from the sample, then the Q value of taking that action should be larger compared to Q values of taking other actions.

4 Experiments

We trained the modules first to get Q_i before running the inverse reinforcement learning algorithm. For each module, the agent only considers the closest target and the closest obstacle. For the path module, the path is defined as segments of waypoints, so the closest waypoint is considered. The agent takes the distance and angle to the closest objects as the state representation.

To make our weights represent the significance of the modules, we normalize the sub-MDPs with the unit (positive or negative) rewards. The reward is 1 for collecting a target, -1 for running into an obstacle. We define the value function directly for the path module to have a path following performance, as it is tricky to give reward for such performance.

We make some constraints on our learning agent to make it walk like a human. We can find in the human trajectories that humans walk smoothly. They don't turn sharply. Our agent is only allowed to do three actions — going straight ahead, turning left with a small step, and turning right with a small step.




Figure 2: The trajectories of humans and our agent in four different tasks. From left to right are Task 1 to Task 4. The active modules should be (Path), (Obstacle + Path), (Target + Path), (Target + Obstacle + Path), respectively.

We report the results in Figure 2, same as Figure 1, the red circles are obstacles. The blue circles are targets. The gray lines are the path. The black lines are trajectories of human — each line represents one human trajectory. Using the weights derived from the algorithm, the trajectories of our agents are drawn in green color.

Average by Task	Num Targs Hit	Num Obst Hit
1	2.34	2.13
2	3.03	0.13
3	10.19	2.28
4	9.88	0.03

Table 1: Number of targets hit and number of obstacles hit of the humans.

Average by Task	Num Targs Hit	Num Obst Hit
1	1.25	1.62
2	3.62	2.37
3	5.14	3.14
4	5.00	2.00

Table 2: Number of targets hit and number of obstacles hit of the learning agent.

Although humans don't agree each other on how to avoid obstacles or collecting targets, our agent can figure out what the humans are doing, and perform a similar trajectory. Table 1 and Table 2 evaluate the performance by showing the number of targets hit and number of obstacles hit. Note that the numbers bold-ed are active modules in the corresponding task. We can observe that humans still do better than our agent in these tasks.




Figure 3: Heatmaps of the log of values of Equation 1 for different weights for the four tasks, respectively. The red zones indicate higher values. The upper two are Task 1 and 2. The lower two are Task 3 and 4.

In Figure 3, we show the log of values of Equation 1 for different weights. We can observe the centroids of red zones move for different tasks. It stays at the origin in Task 1, so none of target and obstacle modules are active. It moves away from the origin when a module is active.

From the heatmaps, we find the optimums exist for these tasks. The optimal weights are also consistent with the experiment context.

5 Conclusion and Future Work

We interpreted human behavior using inverse module reinforcement learning in this paper. We have some positive results, but the performance of our agent is still inferior than humans.

There are also some observations potential for the future work. First, weighted sum of Q functions is one way to combine multiple sub-MDPs. We also propose other ways including, for example, scheduling between different modules, with only one active at one time. This is also called skilled in the literature [2]. However, we adopt the weighted sum approach because this is more reasonable for human behavior. When a human tries to collect targets while avoiding obstacles, these two modules are expected to be both active. A scheduling approach may yield frequent oscillation between these two modules.

Second, we also assumes independency between modules. Correlation between modules doesn't impair our analysis in this paper. In Figure 3, we can tell that the target module and obstacle module tend to be negatively correlated from the shape of the red zones.

Lastly, weights may be dynamic and different from state to state. However, with such assumption, we need to learn a mapping from state to weights. In this case, the curse of dimensionality still exists, and inverse learning would be difficult.

References

- [1] Daphne Koller and Ronald Parr. Computing factored value functions for policies in structured mdps. In *IJCAI*, volume 99, pages 1332–1339, 1999.
- [2] George Konidaris and Andre S Barreto. Skill discovery in continuous reinforcement learning domains using skill chaining. In *Advances in Neural Information Processing Systems*, pages 1015–1023, 2009.
- [3] Constantin A Rothkopf and Dana H Ballard. Modular inverse reinforcement learning for visuomotor behavior. *Biological cybernetics*, 107(4):477–490, 2013.
- [4] Nathan Sprague and Dana Ballard. Multiple-goal reinforcement learning with modular sarsa (0). In *IJCAI*, pages 1445–1447. Citeseer, 2003.
- [5] Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*, volume 1. Cambridge Univ Press, 1998.
- [6] Philip S Thomas and Andrew G Barto. Motor primitive discovery. In *ICDL-EPIROB*, pages 1–8, 2012.