
Modular Inverse Reinforcement Learning on Human Motion

Shun Zhang

Department of Computer Science
University of Texas at Austin
Austin, TX 78712
menie482@cs.utexas.edu

Matthew Tong

Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
mhtong@gmail.com

Mary Hayhoe

Center for Perceptual Systems
University of Texas at Austin
Austin, TX 78712
hayhoe@utexas.edu

Dana Ballard

Department of Computer Science
University of Texas at Austin
Austin, TX 78712
dana@cs.utexas.edu

Abstract

Keywords: Reinforcement learning, human cognition, inverse learning

1 Introduction

Human is able to learn to accomplish complicated tasks much faster than machines can do. However, most of the tasks can be decomposed into subtasks. A human may already have the capacity to accomplish the subtasks. He simply transfers the knowledge of the subtasks and integrate them for the objective task. For example, when a human learns to drive, he has much prior domain knowledge to help him with this objective task. He integrates his skills of controlling the velocity, avoiding other vehicles, navigating, responding to traffic lights, and so on.

Similar ideas are also adopted in the learning literature. Reinforcement learning suffers from the curse of dimensionality. It is inefficient for a learning agent to learn a complicated task from scratch. So learning algorithms including hierarchical reinforcement learning [1], modular reinforcement learning [5] are proposed. They also decompose the objective task into subtasks. The learning algorithms first learn the subtasks independently, then learn how to combine these subtasks.

A natural question to ask is how human integrates skills for subtasks, and whether we can use the same method for integrating subtasks in reinforcement learning problems. In this abstract, we analyze human's behavior of accomplishing a task composite of various subtasks. We collected human motion data when human accomplish a task, and try to interpret the behavior using inverse reinforcement learning. With the best of our knowledge, this approach is novel in the literature.

This abstract is organized as follows. Section 2 introduces the domain of the composite task that we collected human data. Section 3 describes the main algorithm, modular inverse reinforcement learning. We report our experiment results in Section 4, and conclude in Section 5.

2 Multi-objective Sidewalk Domain




Figure 1: The task of collecting targets, avoiding obstacles and following the path. [4]

Consider the task illustrated in Figure 1 A). The avatar is asked to do three sub-tasks simultaneously — 1) following the path, indicated by the gray line on the ground, 2) getting targets, the blue cylinders, and 3) avoiding obstacles, the pink cylinders. This is an experiment design used in the literature to evaluate modular reinforcement learning [4].

From the reinforcement learning perspective, this task can be decomposed to be three sub-tasks as described above. In Figure 1 B), if the agent knows the distance and angle to an object, he is expected to know the optimal action to avoid or pursue it.

3 Modular Reinforcement Learning

Reinforcement learning [6]. Value functions is decomposed [2].

If a human knows the policies of the sub-tasks, or sub-MDPs, he can accomplish a complicated behavior by combining the sub-MDPs. That is,

$$Q(s, a) = \sum_i w_i Q_i(s, a)$$

where Q_i is the Q value of the i-th sub-MDP, w_i is the weight of the i-th sub-MDP. $w_i \geq 0, \sum_i w_i = 1$.

Different weights can yield different performance. Let w_1, w_2, w_3 be weights for the task of target collection, obstacle avoidance, and path following, respectively. Let w be the vector of $(w_i)_1^n$. An agent with $w = (1, 0, 0)$ only collect targets, and one with $w = (0, 0.5, 0.5)$ may avoid the obstacles and follow the path.

To obtain the weights given the samples, we need to use the Inverse Modular Reinforcement Learning technique [4].

4 Experiments

We conducted experiments to ask volunteers to accomplish different tasks and recorded their trajectories. There are four kind of tasks. Task 1, following the path only, and ignoring other objects. Task 2, following the path, while avoid the obstacles. Task 3, following the path, while attain targets when possible. Task 4, following the path, collecting the targets and avoiding obstacles simultaneously. The human data are collected by the Center for Perceptual Systems in University of Texas at Austin.




Figure 2:

From a different perspective, can we find a weight vector to best interpret human’s behavior. In Figure 2, same as Figure 1, the red circles are obstacles. The blue circles are targets. The gray line is the path. The black lines are trajectories of human.

We make some constraints on our learning agent to make it walk like a human. We can find in the human trajectories that humans walk smoothly. They don’t turn sharply. Our agent is allowed to do three actions — going straight ahead, turning left with a small step, and turning right with a small step.

To make our weights represent the significance of the modules, we normalize the sub-MDPs with the unit (positive or negative) rewards. The reward is 1 for collecting a target, -1 for running into an obstacle. We define the value function directly for the path module to have a path following performance.

The agent only considers the closest target and the closest obstacle.

We assume that our learning agent only knows the three sub-MDPs and the human data. It looks at the human behavior, and finds the weights that can interpret such behavior. Using such weights, the trajectories of our agents are drawn in the green lines. We can tell that in the left figure, the agent puts a large weight on path-following. In the right figure, it puts weights on all sub-MDPs.

Average by Task	Num Targs Hit	Num Obst Hit
1	2.34	2.13
2	3.03	0.13
3	10.19	2.28
4	9.88	0.03

Table 1: Number of targets hit and number of obstacles hit of the humans.

Average by Task	Num Targs Hit	Num Obst Hit
1	1.25	1.62
2	3.62	2.37
3	5.14	3.14
4	5.00	2.00

Table 2: Number of targets hit and number of obstacles hit of the learning agent.

In Figure 3, we show the object values for different weights.




Figure 3: Heatmaps of the objective values of different weights for the four tasks, respectively. The red zones indicate higher values. The upper two are Task 1 and 2. The lower two are Task 3 and 4.

5 Conclusion and Future Work

Weighted sum of Q function is one way to combine multiple sub-MDPs. We also propose other ways including, for example, scheduling between different modules, with only one active at one time. This is also called skilled in the literature [3]. However, we adopt the weighted sum approach because this is more reasonable for human behavior. When a human tries to collect targets while avoiding obstacles, these two modules are expected to be both active. A scheduling approach may yield frequent oscillation between these two modules.

We also assume independency between modules. Correlation between modules doesn't impair our analysis in this paper. In Figure 3, we can find that the target module and obstacle module tends to be negatively correlated.

Weights may be dynamic. However, with such assumption, we need to learn a mapping from state to weights. In this case, the curse of dimensionality still exists, and inverse learning would be difficult.

References

- [1] Thomas G Dietterich. Hierarchical reinforcement learning with the maxq value function decomposition. *J. Artif. Intell. Res.(JAIR)*, 13:227–303, 2000.
- [2] Daphne Koller and Ronald Parr. Computing factored value functions for policies in structured mdps. In *IJCAI*, volume 99, pages 1332–1339, 1999.
- [3] George Konidaris and Andre S Barreto. Skill discovery in continuous reinforcement learning domains using skill chaining. In *Advances in Neural Information Processing Systems*, pages 1015–1023, 2009.
- [4] Constantin A Rothkopf and Dana H Ballard. Modular inverse reinforcement learning for visuomotor behavior. *Biological cybernetics*, 107(4):477–490, 2013.
- [5] Nathan Sprague and Dana Ballard. Multiple-goal reinforcement learning with modular sarsa (0). In *IJCAI*, pages 1445–1447. Citeseer, 2003.

- [6] Richard S Sutton and Andrew G Barto. *Reinforcement learning: An introduction*, volume 1. Cambridge Univ Press, 1998.